

# Homework 1 - Design Writeup

Patrick Collins

January 14, 2014

## 1 graph.c

File-level constants:

- Error codes:
  - `INVALID_INPUT_ERROR`  
Used for graphs with invalid distances. Violations of Graph invariant #1.
  - `MALFORMED_INPUT_ERROR`  
Used for files with the wrong number of lines, etc. Violations of Graph invariant #2.
- Values:
  - `INF_DIST` 1000
  - `MAX_DIST` 10000000

### 1.1 Graph

#### 1.1.1 Description

The implementation here closely follows the Wikipedia description of the algorithm linked in the assignment. The description is omitted.

#### 1.1.2 Implementation Overview

Properties:

- `int vertices`
- `int** matrix`

Methods: (Following the Python style of a leading underscore for private methods, providing some degree of information hiding without unnecessarily complicating testing. This convention is followed for the remainder of the document.)

- `int _get_element(Graph* self, int row, int column)`
- `void _set_element(Graph* self, int row, int column, int val)`
- `void _update_dist(Graph* self, int i, int j, int k)`
- `void _update_dists(Graph* self, int k)`
- `bool graph_eq(Graph* g1, Graph* g2)`
- `int get_verticies(Graph* self)`
- `void solve_graph(Graph* self)`
- `void print_graph(Graph* self)`

Constructors:

- `Graph* graph_from_file(file *f)`

Destructors:

- `void free_graph(Graph* g)`

Invariants:

1. For each element  $e \in \text{matrix}$ ,  $e < \text{MAX\_DIST} \parallel e == \text{INF\_DIST}$
2. `rows == columns == verticies`

## 2 tsgraph.c

### 2.1 \_DistArgs

Supporting module for TSGraph.

Properties:

- `TSGraph* tsgraph`
- `int i_start`
- `int i_end`
- `int k`
- `int tid`

Constructor:

- `_DistArgs* _make_dist_args(TSGraph* tsgraph, int i_start, int i_end, int k, int tid)`

Destructor:

- `_free_dist_args(_DistArgs* args)`

## 2.2 TSGraph

### 2.2.1 Implementation Overview

Properties:

- `Graph* graph`
- `int num_threads`
- `int k`
- `volatile bool[] working_threads`

Methods:

- `Graph* get_graph(TSGraph* self)`
- `TSGraph* _update_dist_p(TSGraph* self, int i_start, int i_end, int k, int tid)`
- `void* _update_dist_p_wrapper(_DistArgs* args)`
- `TSGraph* _update_dists_p(TSGraph* self)`
- `bool still_working(TSGraph* self)`
- `int get_k(TSGraph* self)`

Constructor:

- `TSGraph* tsgraph_from_graph(Graph* graph, int num_threads)`

Destructor:

- `void free_tsgraph(TSGraph* self)`

Invariants:

1. No two threads will call `_update_dist_p` at the same time unless the argument `k` is equal in each.

### 2.2.2 Description

Noting the description of the algorithm in terms of its recursive formula (as listed on Wikipedia):

$$\begin{aligned}\text{shortestPath}(i, j, 0) &= w(i, j) \\ \text{shortestPath}(i, j, k + 1) &= \min(\text{shortestPath}(i, j, k), \text{shortestPath}(i, k + 1, k), \\ &\quad \text{shortestPath}(k + 1, j, k))\end{aligned}$$

it is evident that successive evaluations of `shortestPath` are independent for values of  $i$  and  $j$  but not  $k$ . Therefore, the `ThreadsafeGraph` will aim to guarantee that every thread is evaluating `shortestPath` for different values of  $i$  and  $j$  but the same value of  $k$ .

This will be implemented by “chunking” out the work on  $i$ , the outer loop index, and using a `bool` array, `working_threads`, to signal when all threads have completed work. Specifically, for  $n$  threads, `working_threads` will be initialized as an  $n$ -element array containing only `true`. When the  $i$ th thread completes its work, it will set the  $i$ th element of `working_threads` to `false`. When every element of `working_threads` becomes `false`, `self->k` will be incremented, every element of `working_threads` will be reset to `true`, the next loop iteration will continue, until `self->k == self->graph->vertices`.

Based on an implementation of `_update_dist` which will take constant time, `_update_dist_p` will take time that is linear on the difference between `i_min` and `i_max`. For each thread except for the last,

$$i\_max - i\_min == self->graph->vertices / self->num\_threads.$$

The last thread will be within `self->graph->vertices % self->num_threads` of this value. Therefore, work will be divided equally between threads, providing for a good load imbalance.

Using an array of `bool` values for synchronization will ensure that, for every  $i$  and  $j$ , `shortestPath( $i, j, k+1$ )` is evaluated after `shortestPath( $i, j, k$ )`, preventing a data race. Additionally, it avoids starting and joining `self->num_threads` threads per loop iteration, lowering overhead.

## 3 Test Plan

### 3.1 Graph

Hypothesis: Graphs are instantiated correctly.

Tests:

- Every graph is equal to itself.
- The value of each vertex of a newly-initialized graph is equal to the corresponding value in the file it is contained in.
- `INVALID_INPUT_ERROR` is raised for graphs with impossible distances.
- `MALFORMED_INPUT_ERROR` is raised for graphs that do not meet format specifications.

Hypothesis: Graphs are solved correctly.

Tests:

- Nodes with no outgoing edges remain `INF_DIST` from all other nodes after a graph is solved.
- 5 nodes in a horizontal line a unit distance away from each other, bidirectionally, are correctly reported as being 0, 1, 2, 3 and 4 units away from the rightmost node, and vice-versa for the leftmost.
- 5 nodes in a vertical line a unit distance away from each other, bidirectionally, are correctly reported as being 0, 1, 2, 3 and 4 units away from the bottommost node, and vice-versa for the topmost.
- 5 nodes in a horizontal line a unit distance away from each other, bidirectionally, with a 2-unit long path from the leftmost to the rightmost node, are correctly reported as being 0, 1, 2, 3, 3 and 2 units away from each other.
- An undirected graph with two nodes `MAX_DIST - 1` away from a third undergoes no change when solved.

### 3.2 TSGraph

Hypothesis: No two threads will evaluate `shortestPath(i, j, k)` for two different values of  $k$  at the same time. Furthermore,  $k$  is evaluated for each  $i, j$  exactly once.

Tests:

- Given a `TSGraph*` object `tsg` with `tsg->num_threads = n` and `tsg->k = 0`, for  $i = 0$  to  $n - 2$ , evaluate `_update_dist_p(tsg, 0, 0, i)`. Then:
  - `still_working(tsg) == true`.

- `get_k(tsg) == 0`.
- After, evaluate `update_dist_p(tsg, 0, 0, n - 1)`. Then:
  - `still_working(tsg) == true`.
  - `get_k(tsg) == 1`.
- Repeat until `k == tsg->graph->vertices`, which will occur in `k - 1` steps.

Hypothesis: **TSGraph** objects are solved correctly.

Tests:

- Solving the same testcases as listed under the Graph tests yields the same graph as in that case, for 1, 2, 3, 4 and 5 threads.